

```
In [1]: pip install feature-engine
```

```
Collecting feature-engine
  Downloading feature_engine-1.6.2-py2.py3-none-any.whl (328 kB)
----- 328.9/328.9 kB 2.3 MB/s eta 0:00:00
Requirement already satisfied: pandas>=1.0.3 in c:\users\99210\anaconda3\lib\site-packages (from feature-engine) (1.4.4)
Requirement already satisfied: scikit-learn>=1.0.0 in c:\users\99210\anaconda3\lib\site-packages (from feature-engine) (1.0.2)
Requirement already satisfied: numpy>=1.18.2 in c:\users\99210\anaconda3\lib\site-packages (from feature-engine) (1.21.5)
Requirement already satisfied: scipy>=1.4.1 in c:\users\99210\anaconda3\lib\site-packages (from feature-engine) (1.9.1)
Requirement already satisfied: statsmodels>=0.11.1 in c:\users\99210\anaconda3\lib\site-packages (from feature-engine) (0.13.2)
Requirement already satisfied: pytz>=2020.1 in c:\users\99210\anaconda3\lib\site-packages (from pandas>=1.0.3->feature-engine) (2022.1)
Requirement already satisfied: python-dateutil>=2.8.1 in c:\users\99210\anaconda3\lib\site-packages (from pandas>=1.0.3->feature-engine) (2.8.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\99210\anaconda3\lib\site-packages (from scikit-learn>=1.0.0->feature-engine) (2.2.0)
Requirement already satisfied: joblib>=0.11 in c:\users\99210\anaconda3\lib\site-packages (from scikit-learn>=1.0.0->feature-engine) (1.1.0)
Requirement already satisfied: patsy>=0.5.2 in c:\users\99210\anaconda3\lib\site-packages (from statsmodels>=0.11.1->feature-engine) (0.5.2)
Requirement already satisfied: packaging>=21.3 in c:\users\99210\anaconda3\lib\site-packages (from statsmodels>=0.11.1->feature-engine) (21.3)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in c:\users\99210\anaconda3\lib\site-packages (from packaging>=21.3->statsmodels>=0.11.1->feature-engine) (3.0.9)
Requirement already satisfied: six in c:\users\99210\anaconda3\lib\site-packages (from patsy>=0.5.2->statsmodels>=0.11.1->feature-engine) (1.16.0)
Installing collected packages: feature-engine
Successfully installed feature-engine-1.6.2
Note: you may need to restart the kernel to use updated packages.
```

```

In [2]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt

from feature_engine.outliers import Winsorizer
import feature_engine.transformation as vt

# for one hot encoding with sklearn
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier

from collections import Counter
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
%matplotlib inline

# Ignoring Unnecessary warnings
import warnings
warnings.filterwarnings("ignore")

```

```

In [4]: df = pd.read_csv("C:/Users/99210/Downloads/winedataset.csv")

```

```

In [5]: df.head(10)

```

Out[5]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4	5
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4	5
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0	7
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	9.5	7
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5	5

```

In [6]: df.shape

```

Out[6]: (1599, 12)


```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64   
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
In [8]: df.quality.unique()
```

```
Out[8]: array([5, 6, 7, 4, 8, 3], dtype=int64)
```

```
In [9]: plt.figure(1, figsize=(10,10))  
df['quality'].value_counts().plot.pie(autopct="%1.1f%%")  
plt.show()
```



```
In [10]: df.quality.value_counts(ascending=False)
```

```
Out[10]: 5    681  
        6    638  
        7    199  
        4     53  
        8     18  
        3     10  
        Name: quality, dtype: int64
```

```

In [12]: def diagnostic_plots(df, variable, target):
# The function takes a dataframe (df) and
# the variable of interest as arguments.

# Define figure size.
plt.figure(figsize=(20, 4))

# histogram
plt.subplot(1, 4, 1)
sns.histplot(df[variable], bins=30, color = 'r')
plt.title('Histogram')

# scatterplot
plt.subplot(1, 4, 2)
plt.scatter(df[variable], df[target], color = 'g')
plt.title('Scatterplot')

# boxplot
plt.subplot(1, 4, 3)
sns.boxplot(y=df[variable], color = 'b')
plt.title('Boxplot')

# barplot
plt.subplot(1, 4, 4)
sns.barplot(x = target, y = variable, data = df)
plt.title('Barplot')


plt.show()

```


```

In [13]: for variable in df:
diagnostic_plots(df, variable, 'quality')

```



```
In [14]: corr = df.corr()
plt.figure(figsize=(20, 9))
k = 12 #number of variables for heatmap
cols = corr.nlargest(k, 'quality')['quality'].index
cm = np.corrcoef(df[cols].values.T)
sns.set(font_scale=1.25)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_kws={'size':
plt.show()
```



In [15]: `df.isnull().sum()`

```
Out[15]: fixed acidity      0
volatile acidity    0
citric acid         0
residual sugar      0
chlorides           0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                 0
sulphates          0
alcohol            0
quality            0
dtype: int64
```


```
In [17]: def detect_outliers(df, features):
outlier_indices = []

for c in features:
    # 1st quartile
    Q1 = np.percentile(df[c], 25)
    # 3rd quartile
    Q3 = np.percentile(df[c], 75)
    # IQR
    IQR = Q3 - Q1
    # Outlier step
    outlier_step = IQR * 1.5
    # detect outlier and their indeces
    outlier_list_col = df[(df[c] < Q1 - outlier_step) | (df[c] > Q3 + outlier_step)]
    outlier_indices.extend(outlier_list_col)

outlier_indices = Counter(outlier_indices)
multiple_outliers = list(i for i, v in outlier_indices.items() if v > 2)

return multiple_outliers
```

In [18]: `diagnostic_plots(df, 'fixed acidity', 'quality')`



```
In [19]: cols = ['fixed acidity', 'volatile acidity', 'residual sugar',
                 'chlorides', 'free sulfur dioxide', 'total sulfur dioxide',
                 'sulphates', 'alcohol']
```

```
In [20]: lt = vt.LogTransformer(variables = cols)

lt.fit(df)
```

```
Out[20]: LogTransformer(variables=['fixed acidity', 'volatile acidity', 'residual sugar',
                                   'chlorides', 'free sulfur dioxide',
                                   'total sulfur dioxide', 'sulphates', 'alcohol'])
```

```
In [21]: df = lt.transform(df)
```

```
In [22]: bins = (2, 6.5, 8)
group_names = ['bad', 'good']
df['quality'] = pd.cut(df['quality'], bins = bins, labels = group_names)
```

```
In [25]: encoder = LabelEncoder()
```

```
In [24]: df['quality'] = encoder.fit_transform(df['quality'])
```

```
In [26]: df['quality'].value_counts()
```

```
Out[26]: 0    1382
         1     217
         Name: quality, dtype: int64
```

```
In [27]: df.head()
```

```
Out[27]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	2.001480	-0.356675	0.00	0.641854	-2.577022	2.397895	3.526361	0.9978	3.51	-0.579818	2.240710
1	2.054124	-0.127833	0.00	0.955511	-2.322788	3.218876	4.204693	0.9968	3.20	-0.385662	2.282382
2	2.054124	-0.274437	0.04	0.832909	-2.385967	2.708050	3.988984	0.9970	3.26	-0.430783	2.282382
3	2.415914	-1.272966	0.56	0.641854	-2.590267	2.833213	4.094345	0.9980	3.16	-0.544727	2.282382
4	2.001480	-0.356675	0.00	0.641854	-2.577022	2.397895	3.526361	0.9978	3.51	-0.579818	2.240710

```
In [28]: X_train, X_test, y_train, y_test = train_test_split(df.drop('quality', axis=1),
                                                             df['quality'],
                                                             test_size=0.3,
                                                             random_state=0)

X_train.shape, X_test.shape
```

```
Out[28]: ((1119, 11), (480, 11))
```



```
In [29]: # set up the scaler
scaler = StandardScaler()

# fit the scaler to the train set, it will learn the parameters
scaler.fit(X_train)

# transform train and test sets
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [30]: dct = DecisionTreeClassifier(random_state = 42)
svc = SVC(random_state = 42)
rf = RandomForestClassifier(random_state = 42)
logreg = LogisticRegression(random_state = 42)
knn = KNeighborsClassifier()
sgd = SGDClassifier()
```

```
In [31]: classifiers = [
    ("knn" , knn),
    ("rf" , rf),
    ("logreg" , logreg),
    ("svc", svc),
    ("sgd",sgd),
    ("dct",dct)
]
```

```
In [37]: from sklearn.metrics import accuracy_score
```

```
In [38]: for clf_name, clf in classifiers:
    # Fit clf to the training set
    clf.fit(X_train, y_train)

    # Predict y_pred
    y_pred = clf.predict(X_test)
    acc = accuracy_score(y_test, y_pred)

    # Evaluate clf's accuracy on the test set
    print('{:s} score : {:.3f}'.format(clf_name, acc))
```

```
knn score : 0.890
rf score : 0.921
logreg score : 0.906
svc score : 0.912
sgd score : 0.838
dct score : 0.892
```

```
In [39]: param_grid = {
    'n_estimators' : [50,100,200],
    'max_depth': (1,5,10),
    'min_samples_leaf': (1,5,10)
}
```

```
In [40]: gridsearch = GridSearchCV(rf, param_grid=param_grid, scoring='accuracy', cv=5,n_jobs=6)
```

```
In [41]: %%capture
gridsearch.fit(X_train,y_train)
```

```
In [42]: gridsearch.best_params_
```

```
Out[42]: {'max_depth': 10, 'min_samples_leaf': 1, 'n_estimators': 100}
```

```
In [44]: rf = RandomForestClassifier(max_depth= 10, min_samples_leaf = 1, n_estimators = 100)
```

```
In [45]: rf.fit(X_train,y_train)
```

```
Out[45]: RandomForestClassifier(max_depth=10)
```

```
In [46]: pred = rf.predict(X_test)
```

```
In [47]: print("Accuracy Score:",accuracy_score(pred,y_test))
print("classification Report:\n",classification_report(pred,y_test))
```

Accuracy Score: 0.9208333333333333

classification Report:


	precision	recall	f1-score	support
0	0.96	0.95	0.96	434
1	0.58	0.63	0.60	46
accuracy			0.92	480
macro avg	0.77	0.79	0.78	480
weighted avg	0.92	0.92	0.92	480

```
In [48]: cm = confusion_matrix(y_test, pred)

df1 = pd.DataFrame(columns=["Bad", "Good"], index= ["Bad", "Good"], data= cm )

f,ax = plt.subplots(figsize=(4,4))

sns.heatmap(df1, annot=True, cmap="Reds", fmt= '.0f', ax=ax, linewidths = 5, cbar = False)
plt.xlabel("Predicted Label")
plt.xticks(size = 12)
plt.yticks(size = 12, rotation = 0)
plt.ylabel("True Label")
plt.title("Confusion Matrix", size = 12)
plt.show()
```



In [ ]: